



Technische
Universität
Braunschweig



Deep Learning for Remote Sensing

Floodnet Classification

Arjun Arora

Data Science

Outline

- Introduction
- Motivation & Objective
- Post disaster visual scene understanding
- Floodnet dataset & Classification
- Data preparation & preprocessing
- Model training
- Transfer learning
- Experimental setup & Parameters
- Results and predictions
- Conclusion
- Future works and tasks
- References

Introduction

How can computer vision help in Disaster Management?

Rapid floods more common
due to climate change



Fig 1: Pakistan floods drone images

Introduction

How can computer vision help in Disaster Management?

Accurate, timely and understandable information - quick response and recovery on large scale.

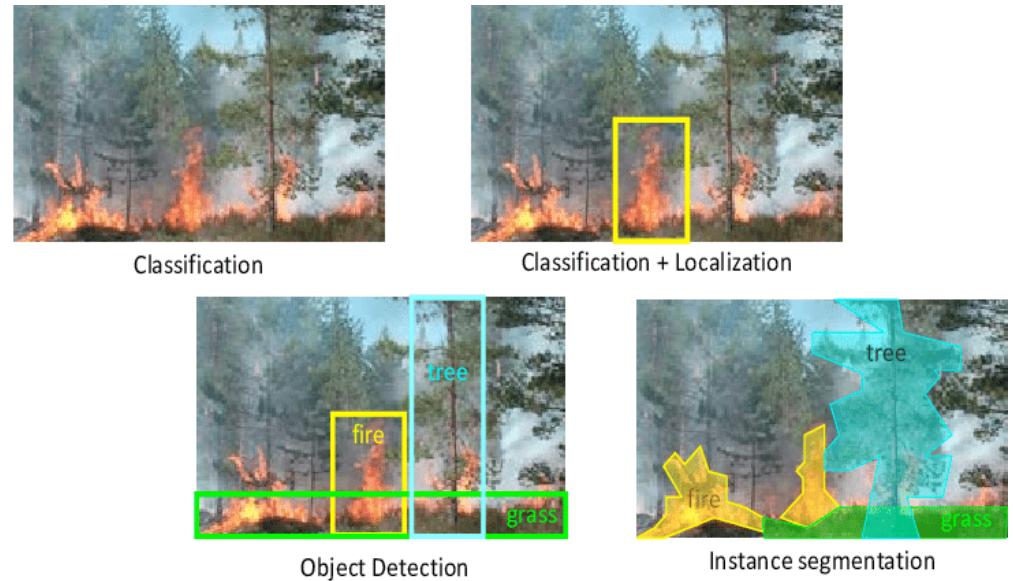


Fig 2: Computer vision in action

Introduction

- Computer vision can help in visual scene understanding post disaster.
- Unmanned Aerial Vehicle(UAV) - access, cheap, HQ images.
- Tasks: Classification, Image Segmentation, Visual Question Answering.
- Deep learning algorithms can analyse the impact of any disaster.
- Understanding of the affected areas and plan the response.

Real Time Disaster Detection System based on UAVs and Deep Learning Algorithms

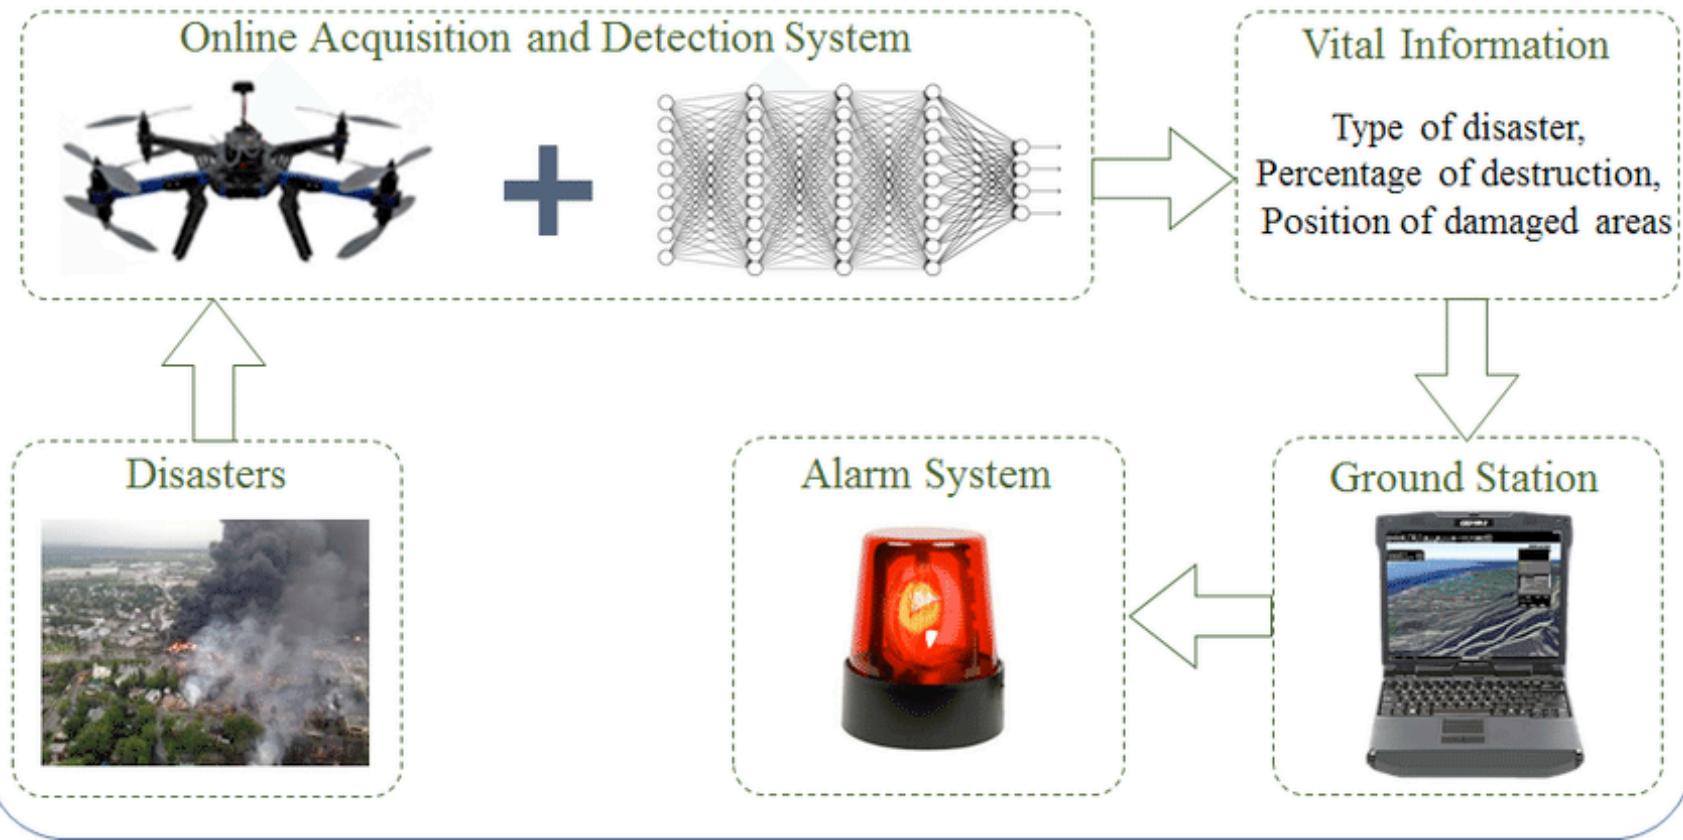


Fig 3: Disaster detection by deep learning

Motivation

- Recent decades, floods are much more common.
- Hurricane Katrina (Florida and Louisiana, August 2005), Hurricane Leslie (France, October 2018), the 2018 monsoon season in India (Kerala, August 2018) and the floods in **Pakistan right now**

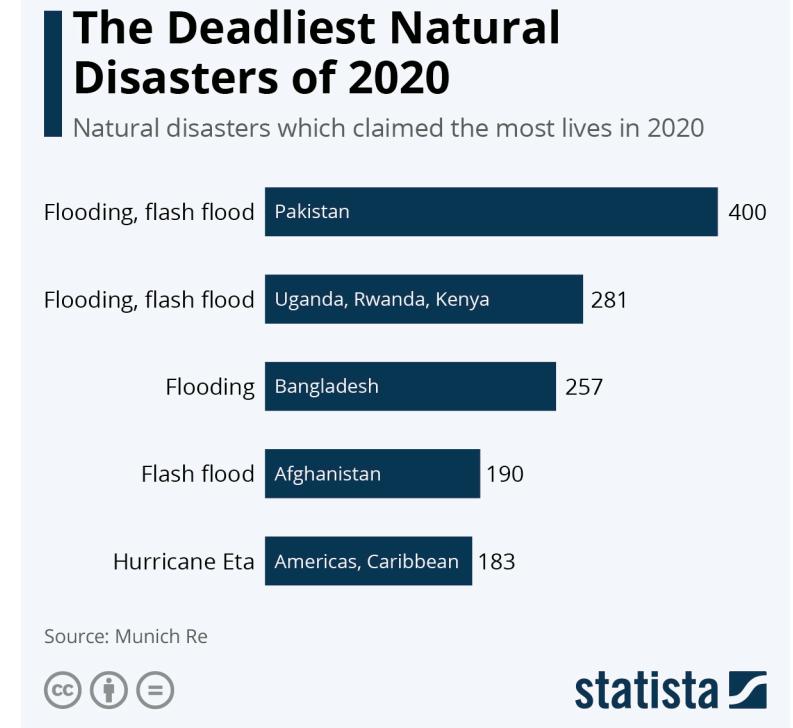


Fig 4: Deadliest natural disasters

Motivation

- Drone is a great tool in mapping, search and rescue, transportation and training.
- Unmanned, easy, safe, HQ images.
- Provides data for quick response in the affected areas.
- NEED: Deep learning models to automatically flag high priority areas post disaster



Fig 5: Flood mapping using drones

Post disaster visual scene understanding

- Classification - Flooded vs Non-Flooded
- Semantic segmentation - flooded roads, buildings, natural water
- Visual question answering -
 - What is the overall condition of the given image?
 - How many buildings are flooded?
 - Condition of the road in this image?

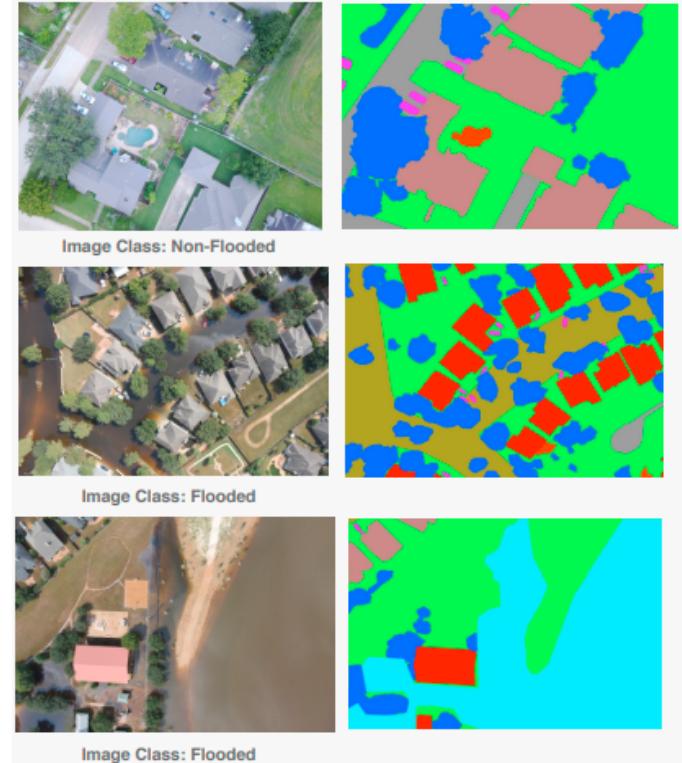


Fig 6: Some examples from Floodnet

Post disaster visual scene understanding- Classification

- Supervised classification - Imagenet, ResNet, EfficientNet
- ResNet - Residual connection architecture; Skip connections solve vanishing gradient.
- EfficientNet - Simple but scalable

Model	Training Accuracy	Test Accuracy	#params
InceptionNetv3	99.03%	84.38%	23.8M
ResNet50	97.37%	93.69%	25.6M
Xception	99.84%	90.62%	22.9M
ResNet18 (our)	96.69%	96.70%	11.6M

Table 1. Classification models comparison

Table 1: Authors classification results

Post disaster visual scene understanding- Semantic segmentation

- Supervised Semantic Segmentation: Extract max info from image.
- Classify each pixel to a class label.
- Models - UNet, PSPNet and DeepLab.

Table 2. Number of images and instances corresponding to different classes.

Object Class	Images	Instances
Building-flooded	275	3573
Building-non-flooded	1272	5373
Road-flooded	335	649
Road-non-flooded	1725	3135
Vehicle	1105	6058
Pool	676	1421
Tree	2507	25889
Water	1262	1784

Table 2: Segmentation classes

Post disaster visual scene understanding- VQA

- VQA framework: Generate questions related to building, road, and entire image.
- Precise damages and situations.
- Find right questions and multiple for each image.
- Challenge: Annotation is difficult



Floodnet Dataset

- Source - small DJI UAV drone after Hurricane Harvey, 2017.
- Images taken during response - realistic and state of practice.
- High resolution images ~ 4/5 Megabytes.
- 1450 train images (2343 total)
- 400 images are labeled.
- Our sample set: Train labeled images
(51 Flooded, 357 Non-flooded)

Floodnet Dataset - Challenges

- Labeled data samples are small.
- Images are very high resolution, difficult to process.
- Resize without loss of efficiency.
- Natural water bodies



Fig 7: Data sample

Floodnet Classification

Classifying hurricane Harvey images as Flooded or Non-Flooded

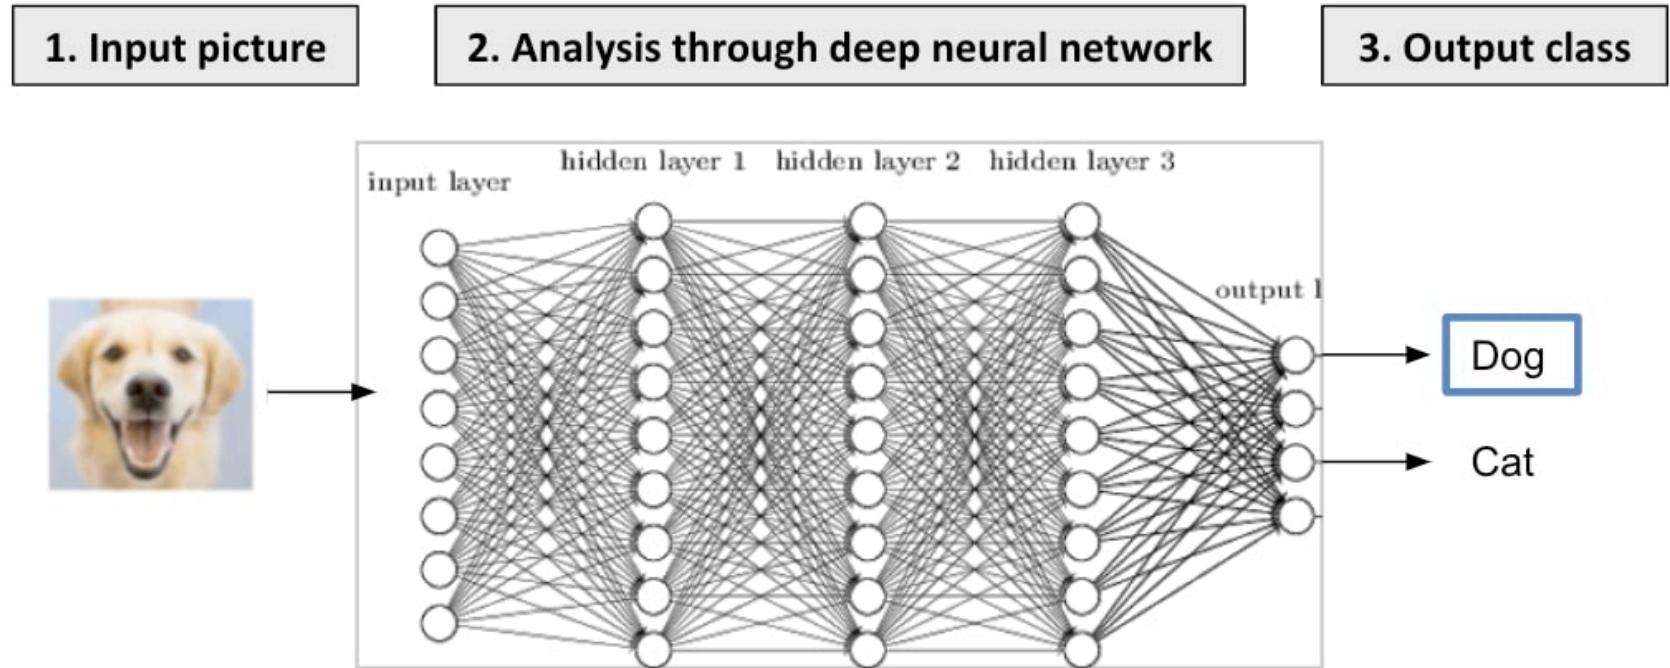


Fig 8: Image classifier

Floodnet Classification

Classifying hurricane Harvey images as Flooded or Non-Flooded

- Steps taken for efficient and accurate classifications:

- Resize the images.
- Split dataset to train, test and validation sets.
- Use ImageDataGenerator - Augmentation and batches.
- Use transfer learning(Resnet 50)
- Test on batches of images.
- Use tensor-board for visualisation and experimentation.

Data preparation & preprocessing

- **Resize the images:** OG image size ~ 7 MBs. Difficult to work with. Large number of parameters working with 4000*3000 pixel images.
- Ram over crashes.
- Solution! Reshape to 400*300

```
# Resize the train flooded and non flooded folders from (4000,3000) to (400,300)
RESIZE=(400,300)
temp_root = "Labeled"
local_root = "output_reshaped"

def resize_and_save(path, resize=RESIZE):
    for img_name in tqdm(os.listdir(os.path.join(temp_root, path))):
        img = cv2.imread(os.path.join(temp_root, path, img_name))
        img = cv2.resize(img, resize)
        cv2.imwrite(os.path.join(local_root, path, img_name), img)
```

Fig 9: Resize the images

Model - ResNet

- Deep residual networks(ResNet-50): 50 layer deep CNN.
- Stacks residual blocks on top of each other to form a network.
- Pretrained on ImageNet database.
- Output layer with sigmoid activation function.

```
from tensorflow.keras.applications import ResNet50

# Initialize the Pretrained Model
feature_extractor = ResNet50(weights='imagenet',
                             input_shape=(400, 300, 3),
                             include_top=False)

# Set this parameter to make sure it's not being trained
feature_extractor.trainable = False

# Set the input layer
input_ = tf.keras.Input(shape=(400, 300, 3))

# Set the feature extractor layer
x = feature_extractor(input_, training=False)

# Set the pooling layer
x = tf.keras.layers.GlobalAveragePooling2D()(x)

# Set the final layer with sigmoid activation function
output_ = tf.keras.layers.Dense(1, activation='sigmoid')(x)

# Create the new model object
model = tf.keras.Model(input_, output_)

# Compile it
model.compile(optimizer='adam',
               loss='binary_crossentropy',
               metrics=['accuracy'])

# Print The Summary of The Model
model.summary()
```

Fig 10: Model preparation

Model - Components

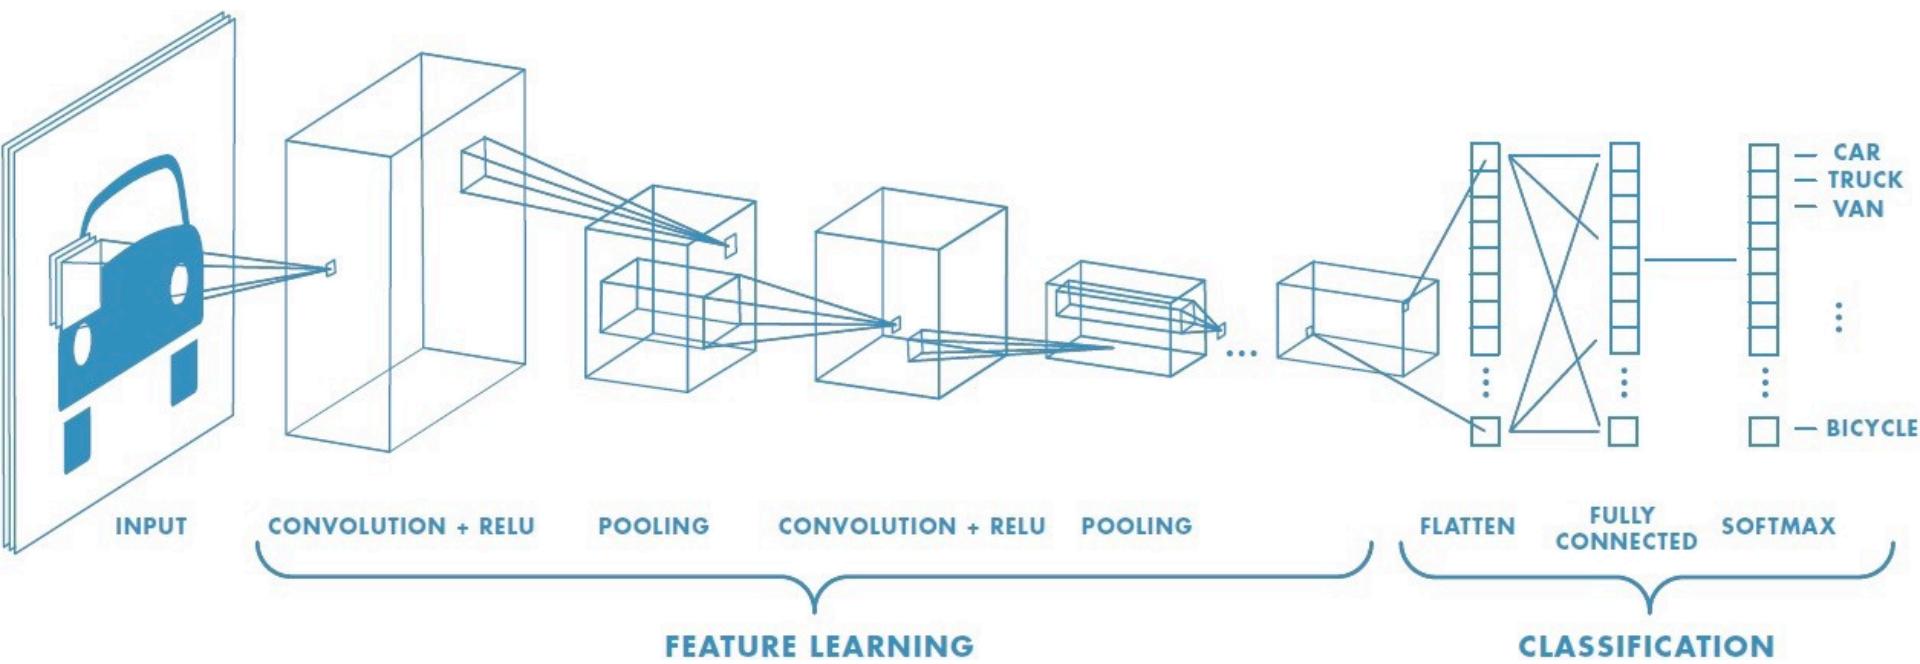
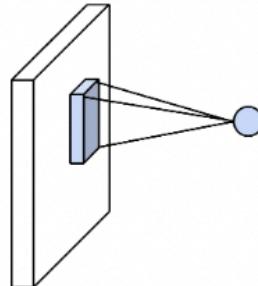


Fig 11: Image classification with conv nets

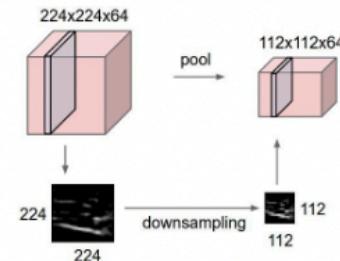
Model - Components

Components of CNNs

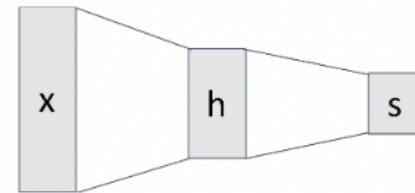
Convolution Layers



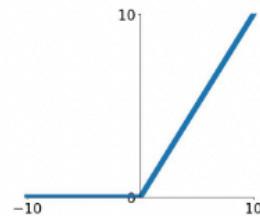
Pooling Layers



Fully-Connected Layers



Activation Function



Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Fig 12: Components of CNN

Model - Architecture

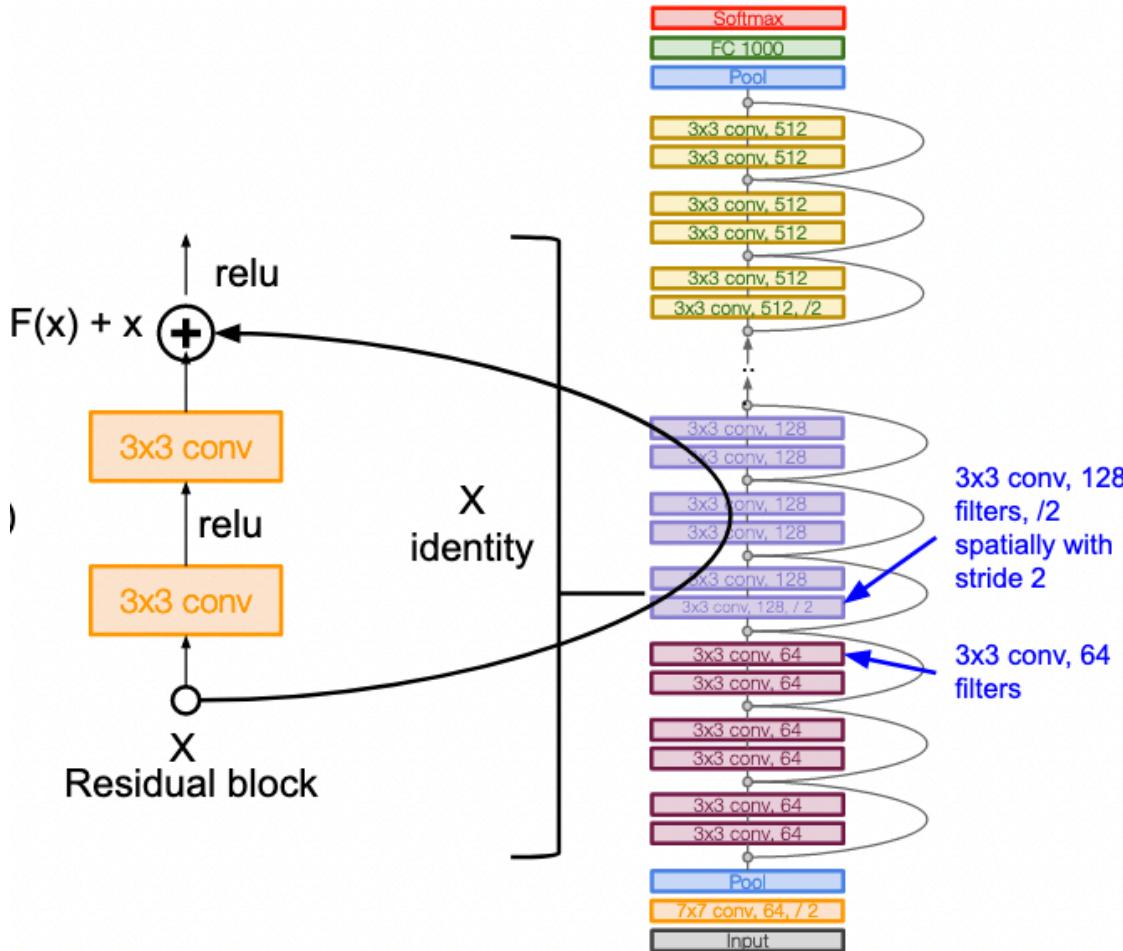


Fig 13: Resnet Architecture

Model - Imagenet

- ImageNet - 14 million images into 1000 object categories.
- Used in visual object detection since 2010.
- A lot of major CNNs have used this dataset to train.
- *These networks also demonstrate a strong ability to generalise to images outside the ImageNet dataset via transfer learning, such as feature extraction and fine-tuning.*



Fig 14 :Imagenet examples

Transfer Learning

Transfer Learning

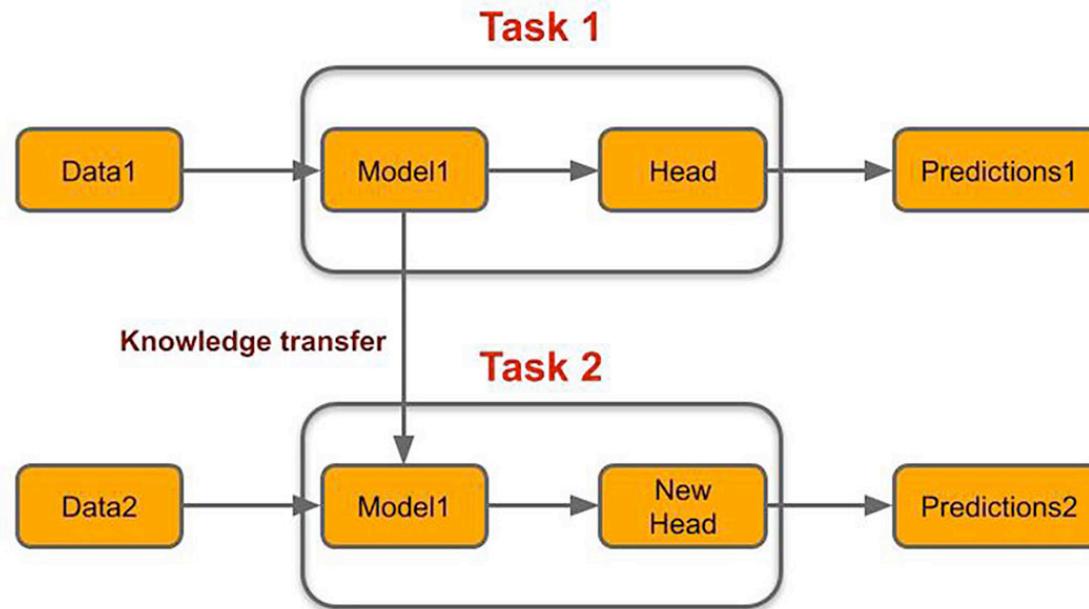


Fig 15: Transfer learning

Transfer Learning

- Reuse pre-trained model on a new problem.
- Little data, quick train times.
- Used in Computer vision and NLP.
- You need to add/modify some layers according to the new task.

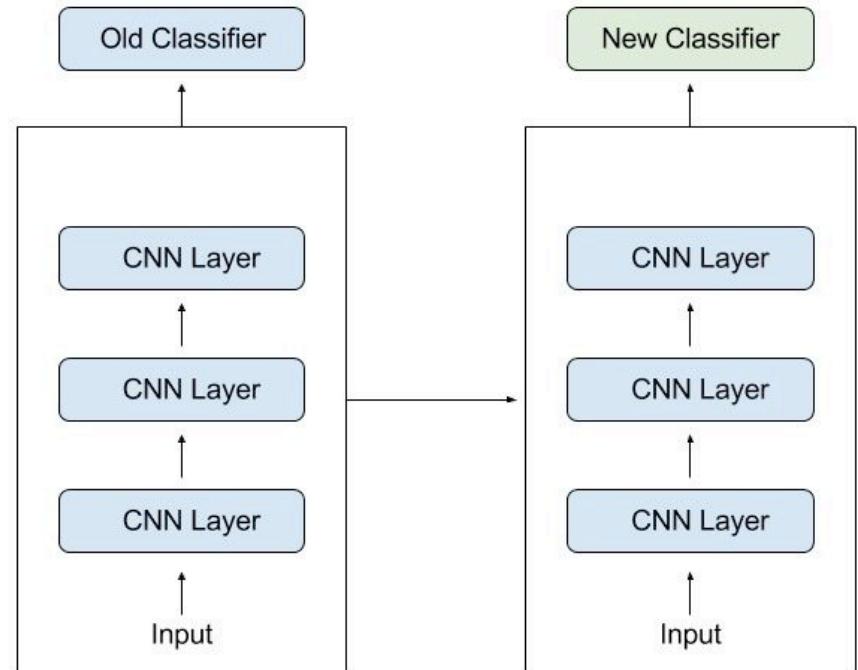


Fig 16: Learning transfer from old task to new task

Experimental setup

- Train - 80%, test - 10%, val - 10%.
- Labels 1 for Flooded, 0 for Non-flooded.
- Input layer + Resnet50 feature extraction layer + Global average pooling layer + Output layer with sigmoid activation.
- **Integrated Feature Extractor:** The pre-trained model, or some portion of the model, is integrated into a new model, but layers of the pre-trained model are frozen during training.

Experimental setup - Model Summary

```
... Model: "model_1"

Layer (type)          Output Shape         Param #
=====
input_4 (InputLayer)   [(None, 400, 300, 3)]  0
resnet50 (Functional) (None, 13, 10, 2048)    23587712
global_average_pooling2d_1 (None, 2048)
(GlobalAveragePooling2D) 0
dense_1 (Dense)        (None, 1)            2049

=====
Total params: 23,589,761
Trainable params: 2,049
Non-trainable params: 23,587,712
```

Fig 17: Model summary



Experimental setup - Layers

- **GlobalAveragePooling layer** - Replace fully connected layer. Take average of each feature map and the result is fed directly to activation.
- **Sigmoid activation function** - Only two classes so like 2-element softmax. Output always between 0 and 1.
- Adam optimiser, binary cross_entropy and metrics.

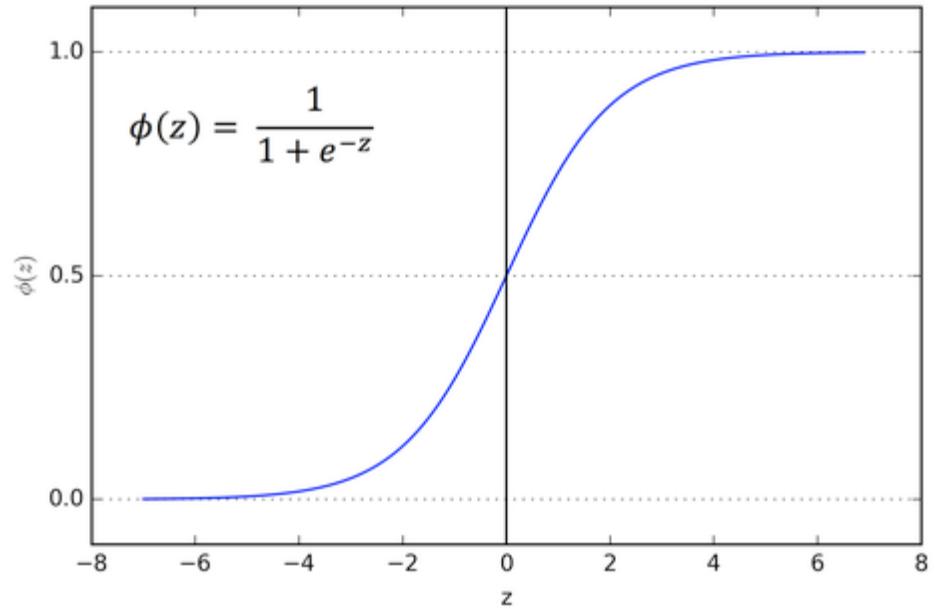


Fig 18: Sigmoid function

Experimental setup - Model compile

- **Optimiser** - Adam optimisation is a stochastic gradient descent method that is based on adaptive estimation of first-order and second-order moments.

- **Loss function** - Binary crossentropy, binary classification tasks.

- **Metrics** - Accuracy. Pred vs actual match.
Total and count. Binary accuracy.

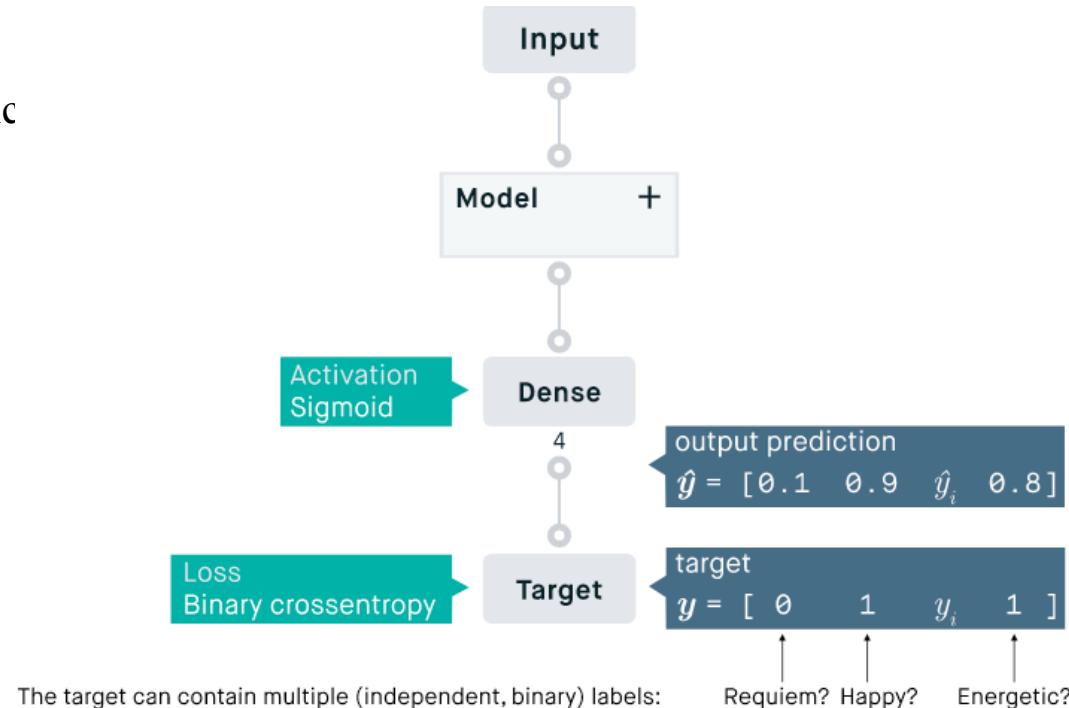


Fig 19: Binary crossentropy

Results - Train & Val

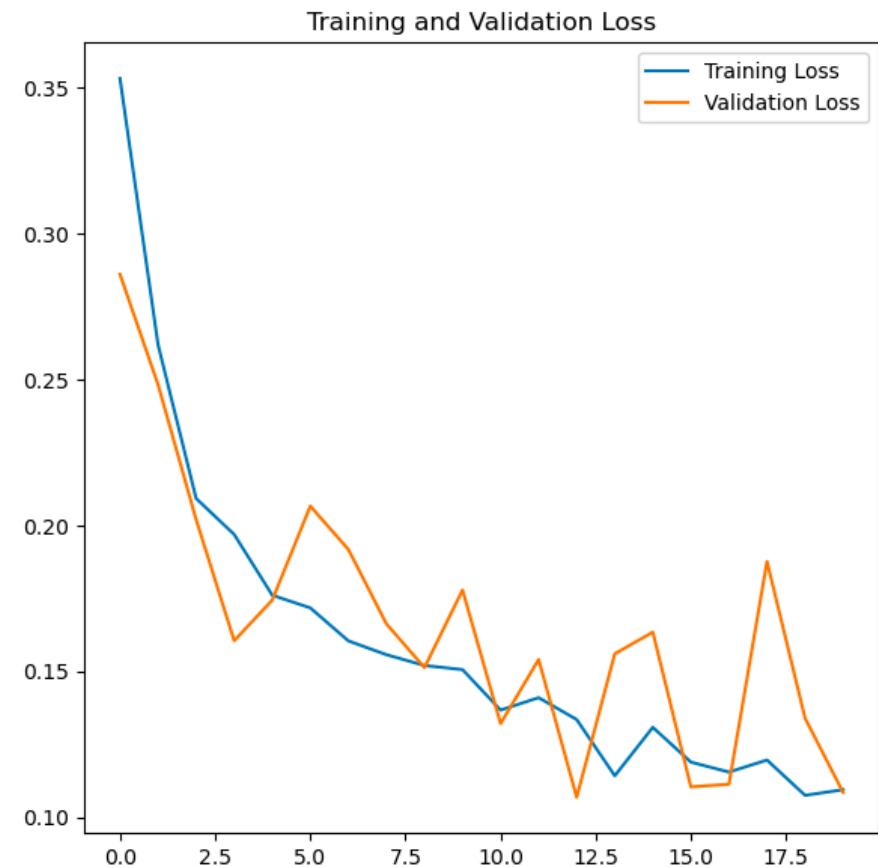
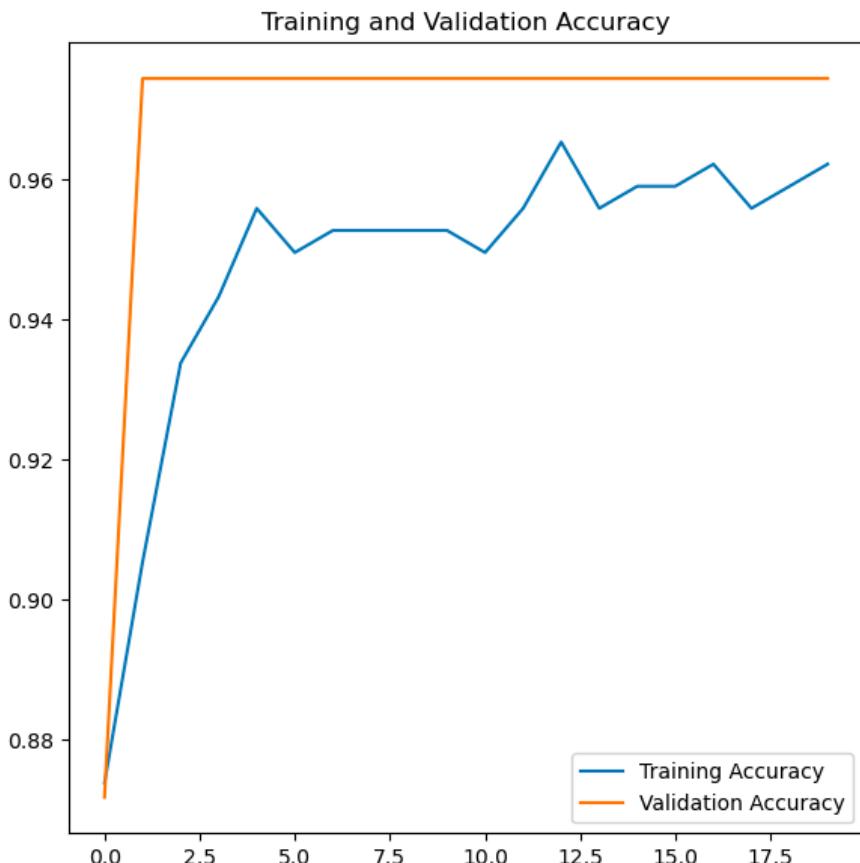


Fig 20: Accuracy and Loss trends

Results - Tensorboard

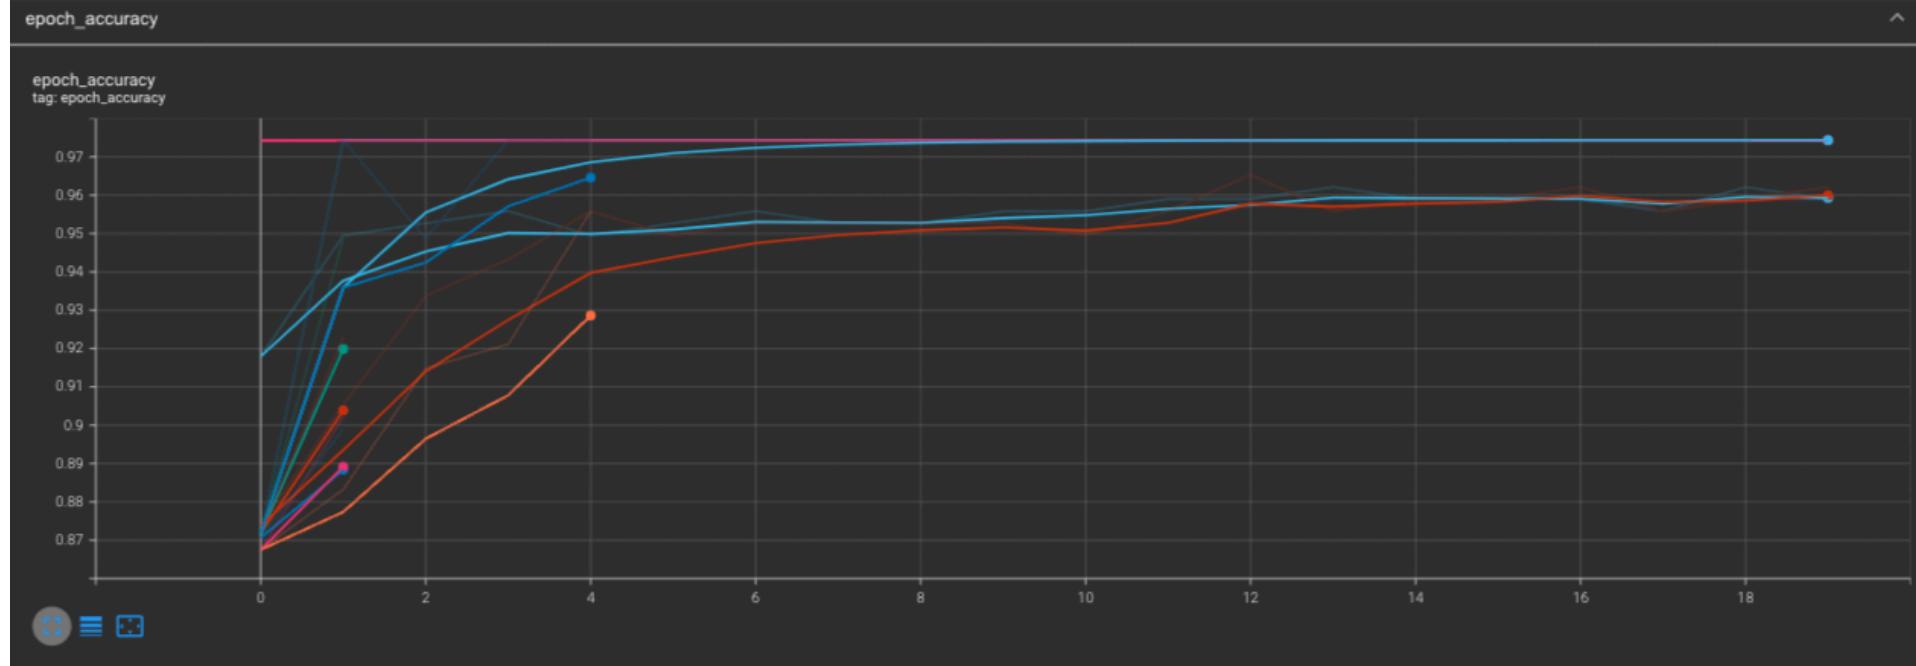


Fig 21: Accuracy trend

Results - Tensorboard

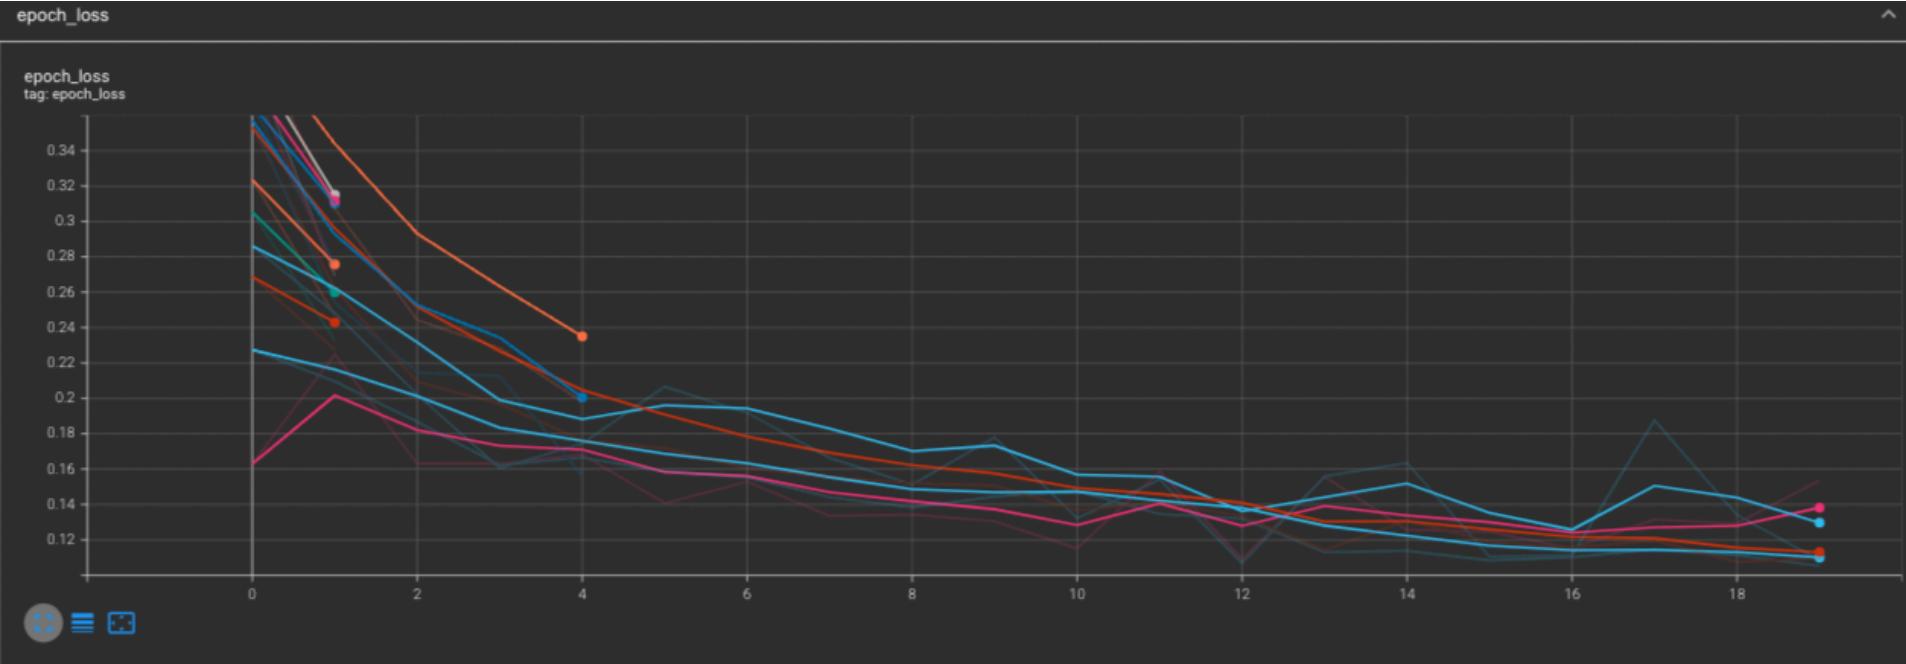


Fig 22: Loss trend

Results - Test data

		precision	recall	f1-score	support
	0	0.89	0.86	0.87	36
	1	0.29	0.33	0.31	6
	accuracy			0.79	42
	macro avg	0.59	0.60	0.59	42
	weighted avg	0.80	0.79	0.79	42
		[[31 5] [4 2]]			

Fig 23: Test accuracy metrics

Results - Examples

```
# predicting images Not flooded class example
img = load_img('output_reshaped/test/10840.jpg', target_size=(img_width, img_height))
x = img_to_array(img)
x = np.expand_dims(x, axis=0)
img.show()

[27]

images = np.vstack([x])
classes = model.predict(images, batch_size=10)
print(classes)

[25]
...
2022-09-15 00:18:00.278551: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_regis
1/1 [=====] - 1s 731ms/step
[[0.06705595]]


df_test.loc[df_test['filename'] == '10840.jpg']

[26]
...
   filename  label
28  10840.jpg      0
```

Fig 24: Non-Flooded example

Results - Examples

```
# predicting images: Yes flooded class
img = load_img('output_reshaped/test/7325.jpg', target_size=(img_width, img_height))
x = img_to_array(img)
x = np.expand_dims(x, axis=0)
img.show()

images = np.vstack([x])
classes = model.predict(images, batch_size=10)
print(classes)

1/1 [=====] - 0s 69ms/step
[[0.90135497]]

df_test.loc[df_test['filename'] == '7325.jpg']
✓ 0.2s

filename  label
0  7325.jpg      1
```

Fig 25: Flooded example

Summary

- Results achieved above 95% accuracy for Train and Validation.
- 80% accuracy for test data.
- Model performed well even for small number of examples using transfer learning.
- Training times low due to reshaping the data.
- Computationally efficient.
- Examples show reliable accuracy in practical settings.

Future works

- Use non-labeled data for training using Semi-supervised techniques.
- Right now we used only around 400 images, total data available > 2000 HQ images.
- Use max info from the images in further tasks like Semantic segmentation which requires pixel-wise details.
- VQA can also be performed and tested against author's results.

References

- <https://www.climatechange.ai/papers/neurips2021/55>
- <https://arxiv.org/pdf/2012.02951.pdf>
- <https://arxiv.org/pdf/2105.08655.pdf>
- <https://github.com/BinaLab/FloodNet-Challenge-EARTHVISION2021>
- <https://github.com/sahilkhose/FloodNet>
- <https://machinelearningmastery.com/how-to-use-transfer-learning-when-developing-convolutional-neural-network-models/>
- <http://cs231n.stanford.edu/schedule.html>